

# FACIAL EMOTION RECOGNITION USING SUPERVISED AND SEMI-SUPERVISED LEARNING

# AGENDA

- INTRODUCTION
- DATASET
- APPROACH & RESULTS
- CONCLUSION
- LIFE IN TOSHIBA AND JAPAN

# INTRODUCTION

- Facial Emotion Recognition Classification (Kaggle challenge 2013)\*

Input image



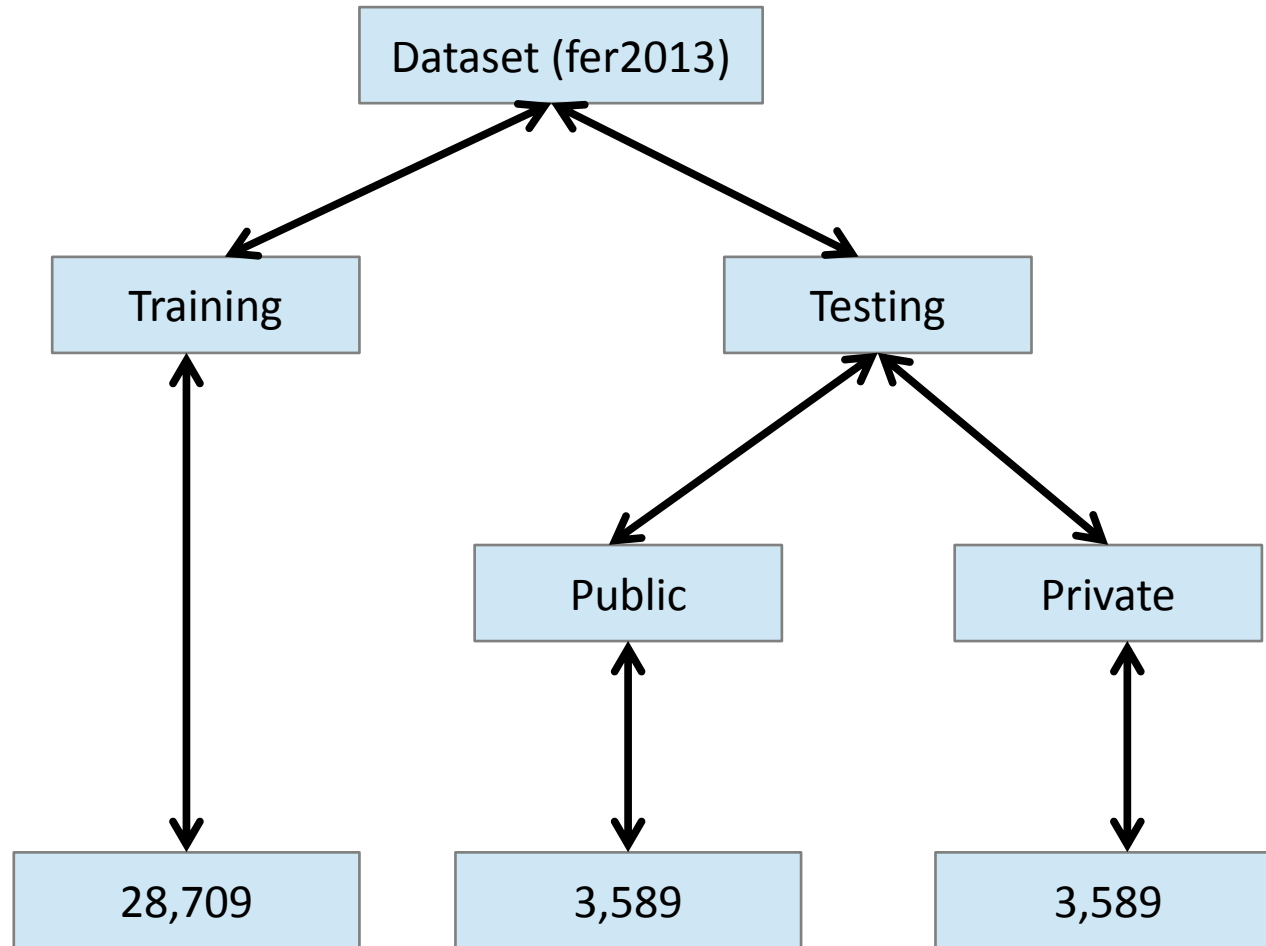
Output category of emotion

0=Angry
1=Disgust
2=Fear
3=Happy
4=Sad
5=Surprise
6=Neutral



\*<https://www.kaggle.com/c/challenges-in-representation-learning-facial-expression-recognition-challenge>

# Dataset



Total images:

[\\*https://www.kaggle.com/c/challenges-in-representation-learning-facial-expression-recognition-challenge](https://www.kaggle.com/c/challenges-in-representation-learning-facial-expression-recognition-challenge)


# Image specification

- Each image is 48 x 48 pixels – 2304 pixel values.
- Images are grayscale in nature – pixel value ranges from 0 to 255.
- Face has been automatically registered – occupies the same amount of space in each image.
- Number of images per class: [3995 436 4097 7215 4830 3171 4965]

Good and clear images



Bad images with thumbnails written across



# APPROACH & RESULTS

Problem of multi-class (our case: 7) categorization.

Many possible approaches:


- k-nearest neighbors (kNN)
- Support Vector Machine (SVM)
- Hierarchical classification (Decision trees)
- Neural networks – Convolutional Neural Networks (CNN)**

Divided into two categories:


1. Fully supervised learning – Training the neural network with all the training images (100%) in the dataset.
2. Semi-supervised learning – Training the neural network with only a part of the training images (10%, 20%) and unlabeled data.

# Fully supervised learning

•Start simple! - With a one layer CNN architecture → accuracy achieved **36%**



•Literature\* suggested to have at least **three CNN layers** for  
•multi-class categorization problem



\*<https://arxiv.org/pdf/1701.08816.pdf>

# Architectural changes and its impact

CNN architectural change	Accuracy / Impact on accuracy (%)
Simple 1 layer CNN	36%
4 conv2D, 3 maxpooling, 2 dense	56%
4 conv2D, 2 maxpooling, 2 dense	56.7% / increased by ~0.5%
4 conv2D, 1 maxpooling, 1 average pooling, 2 dense	58.9% / increased by ~2-2.5%
Adding kernel regularizers with l2 regularization (0.001) in the dense layer	<b>59.7%</b> / increased by ~1.0%
Included strides in the maxpooling layer	58.6% / decreased by ~1.0%
Replace normal dropout with <b>Gaussian dropout</b>	<b>59.7%</b> / almost remains the same
<b>Resizing</b> the image (64 x 64)	<b>59.7%</b> / almost remains the same
Changing the size and number of filters in Conv2D	<b>59.7%</b> / almost remains the same

Architectural changes were inspired from the literature\* of multi-class categorization.

## Main impacts:

- Removing an extra Maxpooling2D layer.
- Replacing the second Maxpooling2D layer with Average pooling layer.
- /2 kernel regularizers.

\*<https://arxiv.org/pdf/1701.08816.pdf>



# Final cnn architecture for fully supervised learning

Layer	Hyperparameters
Convolutional	32 filters, 3 x 3, L2 regularization (0.001), activation: prelu, zeropadding (2, 2)
Convolutional	64 filters, 3 x 3, activation: prelu, zeropadding (1, 1)
Pooling	Maxpool (5 x 5)
Convolutional	128 filters, 3 x 3, activation: prelu, zeropadding (1, 1)
Convolutional	128 filters, 3 x 3, activation: prelu, zeropadding (1, 1)
Pooling	Average pool (3 x 3)
Dense	1024
Dropout	0.25
Dense	1024
Dropout	0.25
Softmax	Fully connected

Total parameters:  
 4,829,319  
Trainable parameters:  
 4,829,255  
Non-trainable parameters:  
 64

CNN model summary		
Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 46, 46, 32)	320
batch_normalization_1 (Batch Normalization)	(None, 46, 46, 32)	128
p_re_lu_1 (PReLU)	(None, 46, 46, 32)	67712
zero_padding2d_1 (ZeroPadding2D)	(None, 50, 50, 32)	0
conv2d_2 (Conv2D)	(None, 48, 48, 64)	18496
p_re_lu_2 (PReLU)	(None, 48, 48, 64)	147456
max_pooling2d_1 (MaxPooling2D)	(None, 9, 9, 64)	0
zero_padding2d_2 (ZeroPadding2D)	(None, 13, 13, 64)	0
conv2d_3 (Conv2D)	(None, 11, 11, 128)	73856
p_re_lu_3 (PReLU)	(None, 11, 11, 128)	15488
zero_padding2d_3 (ZeroPadding2D)	(None, 15, 15, 128)	0
conv2d_4 (Conv2D)	(None, 13, 13, 128)	147584
p_re_lu_4 (PReLU)	(None, 13, 13, 128)	21632
zero_padding2d_4 (ZeroPadding2D)	(None, 17, 17, 128)	0
average_pooling2d_1 (AveragePooling2D)	(None, 5, 5, 128)	0
flatten_1 (Flatten)	(None, 3200)	0
dense_1 (Dense)	(None, 1024)	3277824
p_re_lu_5 (PReLU)	(None, 1024)	1024
dropout_1 (Dropout)	(None, 1024)	0
dense_2 (Dense)	(None, 1024)	1049600
p_re_lu_6 (PReLU)	(None, 1024)	1024
dropout_2 (Dropout)	(None, 1024)	0
dense_3 (Dense)	(None, 7)	7175
activation_1 (Activation)	(None, 7)	0
Total params: 4,829,319		
Trainable params: 4,829,255		
Non-trainable params: 64		

# Data augmentation on the training images

- Instead of training the network with the original images in the training dataset, manipulate the images and then train the network.

- *ImageDataGenerator* in *keras* has this functionality.

Original image



Augmented images



Image rotation, flip, blur

# Impact on accuracy due to synthetic images

Parameters	Accuracy
rotation_range=20, width_shift_range=0.1, height_shift_range=0.1, shear_range=0.1, zoom_range=0.1, horizontal_flip=True	63.14%
rotation_range=40, width_shift_range=0.2, height_shift_range=0.2, shear_range=0.2, zoom_range=0.2, horizontal_flip=True	60.62%
rotation_range=90, width_shift_range=0.4, height_shift_range=0.4, shear_range=0.4, zoom_range=0.4, horizontal_flip=True	48.55%

- Highest accuracy achieved on kaggle competition: 70% ; Approach used > CNNs with **Support Vector Machine** & Transformation to images generating more data.
- Two recent papers claim to achieve accuracy around 75%:
  1. CNN-SIFT algorithm\*
  2. Illumination correction, landmark facial features first#

\*<https://arxiv.org/pdf/1608.02833.pdf>

#<https://arxiv.org/pdf/1612.02903.pdf>

# Confusion matrix

Class	0	1	2	3	4	5	6	Total
0	197	4	44	58	90	8	89	<b>490</b>
1	17	8	10	9	9	1	1	<b>55</b>
2	47	2	153	46	145	59	76	<b>528</b>
3	13	0	14	765	33	12	42	<b>879</b>
4	41	0	43	63	298	5	144	<b>594</b>
5	10	0	63	45	20	258	20	<b>416</b>
6	17	1	29	66	107	12	394	<b>626</b>
<b>Total</b>	<b>342</b>	<b>15</b>	<b>356</b>	<b>1052</b>	<b>702</b>	<b>355</b>	<b>766</b>	<b>3588</b>

# Semi-supervised learning

- Selecting the ratio of labeled / unlabeled data.

Percentage (labeled/unlabeled)	Accuracy (%) of the built CNN architecture
(100/0)	59
(50/50)	53
(30/70)	49
<b>(20/80)</b>	<b>45</b>
<b>(10/90)</b>	<b>40</b>

- Selecting the architecture for semi-supervised learning with
- Neural Networks:
  - 1. Entropy minimization<sup>\*</sup>
  - 2. Autoencoder<sup>#</sup>
  - 3. **Ladder Network**<sup>\$</sup>
  - 3. Mean Teacher<sup>β</sup>

<sup>\*</sup> <https://papers.nips.cc/paper/2740-semi-supervised-learning-by-entropy-minimization>

<sup>#</sup> <http://proceedings.mlr.press/v27/baldi12a/baldi12a.pdf>

<sup>\$</sup> <https://arxiv.org/abs/1507.02672>

<sup>β</sup> <https://arxiv.org/abs/1703.01780>

# Ladder network




Image source: Main paper (Ladder Network)

### Main parts:

1. Encoder – clean and noisy
2. Decoder
3. Denoising function

### Loss function:

1. **Categorical cross entropy** - Noisy Y and labeled data.
2. **Mean squared error** – Between the outputs of the decoder and clean encoder at each layer.

### Implementation:

- Encoder - **Pre-trained** CNN architecture designed for supervised learning with only 10% labeled data.
- Decoder - Basically an inverse of an encoder - **Deconvolutional** and **upsampling**.
- Denoising function - Function which **reconstructs** a noisy signal from encoder to a clean signal which will be compared with the clean encoder.

# Breaking the ladder network architecture

**Algorithm 1** Calculation of the output  $\mathbf{y}$  and cost function  $C$  of the Ladder network

```

Require:  $\mathbf{x}(n)$ 
# Corrupted encoder and classifier
 $\tilde{\mathbf{h}}^{(0)} \leftarrow \tilde{\mathbf{z}}^{(0)} \leftarrow \mathbf{x}(n) + \text{noise}$ 
for  $l = 1$  to  $L$  do
   $\tilde{\mathbf{z}}^{(l)} \leftarrow \text{batchnorm}(\mathbf{W}^{(l)} \tilde{\mathbf{h}}^{(l-1)}) + \text{noise}$ 
   $\tilde{\mathbf{h}}^{(l)} \leftarrow \text{activation}(\gamma^{(l)} \odot (\tilde{\mathbf{z}}^{(l)} + \beta^{(l)}))$ 
end for
 $P(\tilde{\mathbf{y}} | \mathbf{x}) \leftarrow \tilde{\mathbf{h}}^{(L)}$ 
# Clean encoder (for denoising targets)
 $\mathbf{h}^{(0)} \leftarrow \mathbf{z}^{(0)} \leftarrow \mathbf{x}(n)$ 
for  $l = 1$  to  $L$  do
   $\mathbf{z}_{\text{pre}}^{(l)} \leftarrow \mathbf{W}^{(l)} \mathbf{h}^{(l-1)}$ 
   $\mu^{(l)} \leftarrow \text{batchmean}(\mathbf{z}_{\text{pre}}^{(l)})$ 
   $\sigma^{(l)} \leftarrow \text{batchstd}(\mathbf{z}_{\text{pre}}^{(l)})$ 
   $\mathbf{z}^{(l)} \leftarrow \text{batchnorm}(\mathbf{z}_{\text{pre}}^{(l)})$ 
   $\mathbf{h}^{(l)} \leftarrow \text{activation}(\gamma^{(l)} \odot (\mathbf{z}^{(l)} + \beta^{(l)}))$ 
end for

# Final classification:
 $P(\mathbf{y} | \mathbf{x}) \leftarrow \mathbf{h}^{(L)}$ 
# Decoder and denoising
for  $l = L$  to  $0$  do
  if  $l = L$  then
     $\mathbf{u}^{(l)} \leftarrow \text{batchnorm}(\tilde{\mathbf{h}}^{(L)})$ 
  else
     $\mathbf{u}^{(l)} \leftarrow \text{batchnorm}(\mathbf{V}^{(l+1)} \hat{\mathbf{z}}^{(l+1)})$ 
  end if
   $\forall i: \hat{z}_i^{(l)} \leftarrow g(\tilde{z}_i^{(l)}, u_i^{(l)})$  # Eq. (2)
   $\forall i: \hat{z}_{i,\text{BN}}^{(l)} \leftarrow \frac{\hat{z}_i^{(l)} - \mu_i^{(l)}}{\sigma_i^{(l)}}$ 
end for

# Cost function  $C$  for training:
 $C \leftarrow 0$ 
if  $t(n)$  then
   $C \leftarrow -\log P(\tilde{\mathbf{y}} = t(n) | \mathbf{x}(n))$ 
end if
 $C \leftarrow C + \sum_{l=0}^L \lambda_l \left\| \mathbf{z}^{(l)} - \hat{\mathbf{z}}_{\text{BN}}^{(l)} \right\|^2$  # Eq. (3)
    
```

Image source: Main paper (Ladder Network)

Noisy encoder and clean encoder:

$\mathbf{x}_{\text{n}}, \mathbf{z}_{\text{n}}^{(1)}, \dots, \mathbf{z}_{\text{n}}^{(L)}, \mathbf{y}_{\text{n}} = \text{Encoder}_{\text{noisy}}$   
 $\mathbf{x}, \mathbf{z}^{(1)}, \dots, \mathbf{z}^{(L)}, \mathbf{y} = \text{Encoder}_{\text{clean}}$   
 $\mathbf{x}^{\wedge}, \mathbf{z}^{\wedge(1)}, \dots, \mathbf{z}^{\wedge(L)} = \text{Decoder}$

Encoder [additional parameters before activation]:

**Batchnormalization()**  
**Mean and scaling ()**

Noisy encoder [additional parameters to the clean encoder]

**Gaussian noise ( $\sigma = 0.2$ )** added to each layer.

Loss function:

1. Supervised learning - Categorical cross entropy (element wise multiplication, passed to the below layer)
2. Unsupervised learning – Mean squared error ( $\mathbf{z}^{(L)}, \mathbf{z}^{\wedge(L)}$ )

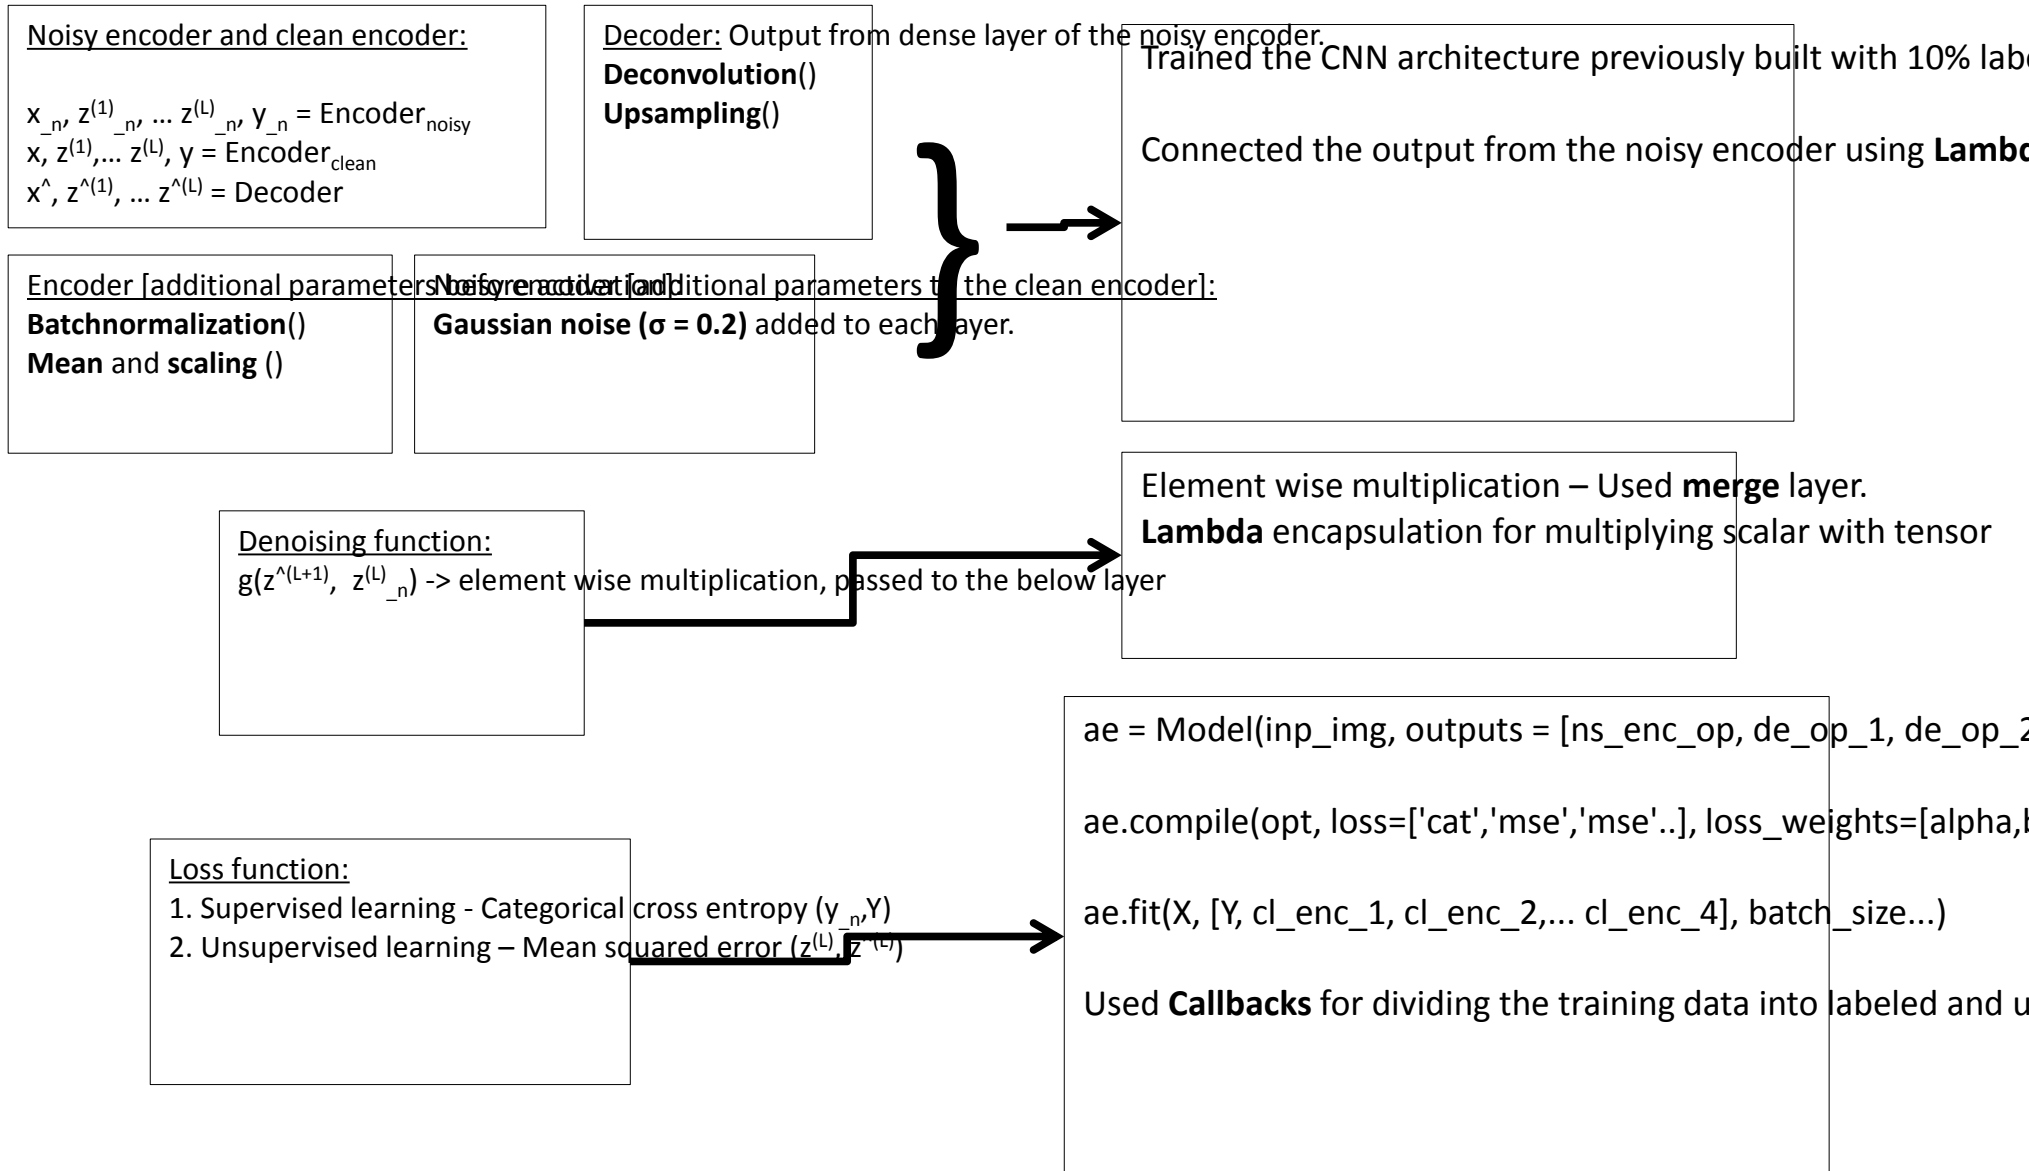
Denoising function:

$g(\hat{z}_i^{(l)}, u_i^{(l)}) = \hat{z}_i^{(l)} + \sigma_i^{(l)} \cdot u_i^{(l)}$

Decoder: Output from dense layer of the noisy encoder

**Deconvolution()**  
**Upsampling()**

# Implementation of the ladder network architecture





# Preliminary Result

- Testing with **10%** labeled data:

- Baseline accuracy of supervised learning – 40%

- With semi-supervised learning – no denoising function yet – 43.5%

- Future work:

- > Test with different variations of the Gaussian noise ( $\sigma = 0.1, 0.3, 0.5$ ).

- > Test with variations in architecture. Only one layer denoising function.

- > Test with different denoising functions.

- > Testing with different ratio of labeled to unlabeled data.

# CONCLUSION

- Dataset needs to be studied properly:

- > **Manipulating** the dataset features which may suit the CNN architecture.

- > Illumination and occlusion inclusion.

- > Class **balance** in the dataset.

- Supervised learning can achieve **exceptional accuracy**, however, **labeled data** for many categorization tasks is not readily available. At the same time, it does not reflect how humans learn.

- Semi-supervised learning:

- > Can **improve** the accuracy compared to only supervised learning with a very **small amount** of labeled data.

- > Diverse applications in many fields where there is a scarcity of labeled data.

- > Combine CNN with some **clustering** technique for predicting the inherent structure in the dataset.

- > More intuitive and related to how humans learn.

- Using *keras*:

- > Suitable for simple model implementation and initial start-up.

- > Modifications in layer outputs, going deep in the network, changing the architecture, combining two different models – better switch to Tensorflow, PyTorch.